

Human Protein Recognition and Predicting Home Values

Mike Crumley

Project Overviews

① Human Protein Recognition




- ▶ Recognize presence/non-presence of 28 protein structures from microscopic cell images
- ▶ Multi-class, multi-label classification
- ▶ Software: Keras (R, on top of TensorFlow)
- ▶ Metric: Macro F_1 score

② Predicting House Prices


- ▶ Predict the sale price of a home using 78 predictors
- ▶ Regression
- ▶ Software: R (various packages)
- ▶ Metric: root mean square logarithmic error

Protein Recognition: Data Exploration


- Training data: 31,072 RGB images (512×512)
- Test data: 11,702 images



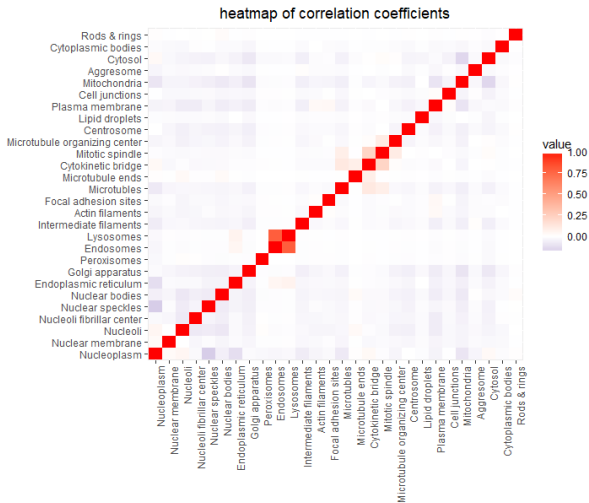
Protein Recognition: Data Exploration




Protein Recognition: Data Exploration



Protein Recognition: Data Exploration



Protein Recognition: Data Exploration



Protein Recognition: Metric

For a single protein structure:

$$F_1 = \left(\frac{\text{precision}^{-1} + \text{recall}^{-1}}{2} \right)^{-1}$$

precision = fraction of images predicted to contain the structure
which actually contained it

recall = fraction of images containing the structure
which were predicted to be so

For the entire multi-class problem:

$$\overline{F}_1 = \left(\frac{\overline{\text{precision}}^{-1} + \overline{\text{recall}}^{-1}}{2} \right)^{-1}$$

Protein Recognition: Metric

Important notes about F_1 :

- Penalizes false negatives more than false positives
- A baseline classifier for F_1 predicts *all* structures to be present in *all* images. A baseline macro F_1 score for this project is

$$F_{1,B} = .1103$$

Baseline scores for individual structures depend heavily on their prevalence; they range from .0007 to .586.

- It is non-differentiable, and can't be used as a loss function (though something close to it can be).

Protein Recognition: Convolutional Networks

Convolutional networks are the thing to do for image recognition.

An example of a convolutional neural network in Keras:

```
model <- keras_model_sequential() %>%  
  layer_conv_2d(filters = 16, kernel_size = c(10, 10), activation = "relu",  
                input_shape = c(512, 512, 3)) %>%  
  layer_max_pooling_2d(pool_size = c(5, 5)) %>%  
  layer_conv_2d(filters = 32, kernel_size = c(10, 10), activation = "relu") %>%  
  layer_max_pooling_2d(pool_size = c(4, 4)) %>%  
  layer_conv_2d(filters = 64, kernel_size = c(3, 3), activation = "relu") %>%  
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%  
  layer_conv_2d(filters = 128, kernel_size = c(3, 3), activation = "relu") %>%  
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%  
  layer_flatten() %>%  
  layer_dense(units = 256, activation = "relu") %>%  
  layer_dense(units = 1, activation = "sigmoid")
```

Protein Recognition: Convolutional Networks

```
## Model
## -----
## Layer (type)                Output Shape                Param #
## -----
## conv2d_1 (Conv2D)           (None, 503, 503, 16)        4816
## -----
## max_pooling2d_1 (MaxPooling2D) (None, 100, 100, 16)        0
## -----
## conv2d_2 (Conv2D)           (None, 91, 91, 32)          51232
## -----
## max_pooling2d_2 (MaxPooling2D) (None, 22, 22, 32)          0
## -----
## conv2d_3 (Conv2D)           (None, 20, 20, 64)          18496
## -----
## max_pooling2d_3 (MaxPooling2D) (None, 10, 10, 64)          0
## -----
## conv2d_4 (Conv2D)           (None, 8, 8, 128)           73856
## -----
## max_pooling2d_4 (MaxPooling2D) (None, 4, 4, 128)           0
## -----
## flatten_1 (Flatten)         (None, 2048)                 0
## -----
## dense_1 (Dense)              (None, 256)                  524544
## -----
## dense_2 (Dense)              (None, 1)                    257
## =====
## Total params: 673,201
## Trainable params: 673,201
## Non-trainable params: 0
## -----
```

Protein Recognition: Special Considerations

- Highly standardized laboratory images
- Extreme rarity of some structures
- Poor correlation between structures
- Typical loss functions (binary crossentropy, categorical cross entropy) can be very poorly aligned with the F_1 metric, especially for rare structures
- Validation/testing will suffer from rarity. Beware of over-fitting to validation.

Protein Recognition: Two Approaches

- ① Training 28 individual binary classification models and ensembling the results. Advantages:
 - ▶ Can artificially balance individual classes, allowing networks to see more examples of rare structures
 - ▶ Low correlation between structures partially justifies this
 - ▶ Can tailor network architecture/optimizer for each individual structure
 - ▶ Possibly lower network complexity (my GPU is not impressive)
 - ▶ Learn more
- ② Training a single model to predict all 28 protein structures simultaneously. Advantages:
 - ▶ Easier and quicker
 - ▶ Can use a custom loss to train directly for macro F_1
 - ▶ Low-mid level representations learned by a network may be useful for many structures

Protein Recognition: 28 Separate Models

Possible tools to combat rarity of structures:

- Custom loss function (penalize false negatives more than false positives)
- Data augmentation (geometric transformations to generate more training images)
- Artificial class balancing (allow network to see a higher proportion of structure-containing images)

Protein Recognition: 28 Separate Models

```
# custom loss/metrics -----

f1.loss = function(y_true, y_pred)
{
  tp = k_mean(y_true*y_pred)
  fp = k_mean((1-y_true)*y_pred)
  fn = k_mean(y_true*(1-y_pred))
  f1 = 2*tp/(2*tp+fn+fp+k_epsilon())
  return((1-f1))
}

F1_macro.fcn.hard = function(y_true, y_pred)
{
  y_pred_hard = tf$floor(2*y_pred/(1+k_epsilon()))
  tp = tf$diag_part(k_dot(k_transpose(y_true), y_pred_hard))
  fn = tf$diag_part(k_dot(k_transpose(y_true), 1-y_pred_hard))
  fp = tf$diag_part(k_dot(k_transpose(1-y_true), y_pred_hard))
  prec = tp/(tp+fp+k_epsilon())
  rec = tp/(tp+fn+k_epsilon())
  F1 = 2*(k_mean(prec)*k_mean(rec))/(k_mean(prec)+k_mean(rec)+k_epsilon())
  return(F1)
}

F1_macro_metric_hard <- custom_metric("Hard_F1", F1_macro.fcn.hard)
```

Protein Recognition: 28 Separate Models

image data generators for training -----

```
train_datagen_rare = image_data_generator(  
    featurewise_center = FALSE,  
    samplewise_center = FALSE,  
    featurewise_std_normalization = FALSE,  
    samplewise_std_normalization = FALSE,  
    zca_whitening = FALSE,  
    zca_epsilon = 1e-06,  
    rotation_range = 45,  
    width_shift_range = .2,  
    height_shift_range = .2,  
    brightness_range = NULL,  
    shear_range = 0,  
    zoom_range = 0,  
    channel_shift_range = 0,  
    horizontal_flip = TRUE,  
    vertical_flip = TRUE,  
    rescale = 1/255,  
)  
  
train_datagen_notrare = image_data_generator(  
    rescale = 1/255,  
    rotation_range = 25,  
    width_shift_range = .1,  
    height_shift_range = .1,  
    shear_range = 0,  
    zoom_range = 0,  
    horizontal_flip = TRUE,  
    vertical_flip = TRUE,  
)
```


Protein Recognition: 28 Separate Models

```
# image data generators for validation -----
```

```
validate_datagen_rare = image_data_generator(  
    featurewise_center = FALSE,  
    samplewise_center = FALSE,  
    featurewise_std_normalization = FALSE,  
    samplewise_std_normalization = FALSE,  
    zca_whitening = FALSE,  
    zca_epsilon = 1e-06,  
    rotation_range = 10,  
    width_shift_range = .05,  
    height_shift_range = .05,  
    brightness_range = NULL,  
    shear_range = 0,  
    zoom_range = 0,  
    channel_shift_range = 0,  
    horizontal_flip = TRUE,  
    vertical_flip = TRUE,  
    rescale = 1/255  
)  
  
validate_datagen_notrare = image_data_generator(  
    rescale = 1/255,  
    horizontal_flip = TRUE,  
    vertical_flip = TRUE  
)
```

Protein Recognition: 28 Separate Models

```
# image data generators for testing -----
```

```
test_datagen_rare = image_data_generator(  
    horizontal_flip = True,  
    vertical_flip = True,  
    rescale = 1/255  
)
```

```
test_datagen_notrare = image_data_generator(  
    rescale = 1/255  
)
```

Protein Recognition: 28 Separate Models

```
# custom training generator for artificial class balancing -----
```

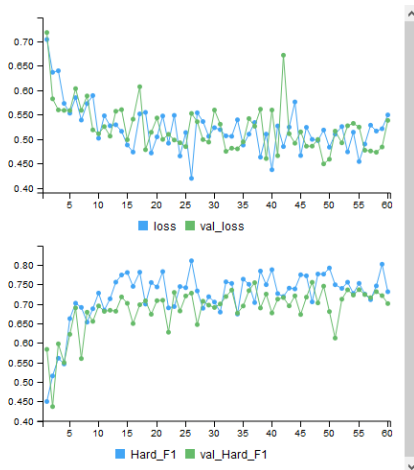
```
biased_rgb_train_generator = function()
{
  function()
  {
    rand.numbers = runif(batch_size)
    present_indices = which(rand.numbers < training.bias[structure])
    notpresent_indices = setdiff(1:batch_size, present_indices)
    training_array = array(dim = c(batch_size, resolution, 3))
    label_array = array(dim = c(batch_size), 0)
    for (i in present_indices)
    {
      training_array[i, , , ] = generator_next(train_generator_present)
    }
    for (i in notpresent_indices)
    {
      training_array[i, , , ] = generator_next(train_generator_notpresent)
    }
    label_array[present_indices] = 1
    return(list(training_array, label_array))
  }
}
```

Protein Recognition: 28 Separate Models

```
# train a bunch of models on 28 structures -----
```

```
train.existing.models = function(  
  model.folder.name,  
  structures = protein.names,  
  resolution = c(512,512),  
  training.bias = default.training.bias,  
  batch_size = 20,  
  steps_per_epoch = 20,  
  epochs = 20,  
  save.history.name = "history.Rda",  
  validation_steps = 100)  
{  
  ...  
}
```

Protein Recognition: 28 Separate Models



Protein Recognition: 28 Separate Models

```
# score a trained model on test data, summarize, find optimal thresholds -----
```

```
F1.score = function(  
  model.folder.name,  
  structures = protein.names,  
  num_present_to_score = 1000,  
  num_notpresent_to_score = 1000,  
  resolution = c(512,512),  
  save.name = "F1_summary.Rda",  
  adjust.thresholds = TRUE)  
{  
  ...  
}
```

Protein Recognition: 28 Separate Models

```
# find best models for each structure -----  
  
find.best.ensemble(  
  model.folder.names,  
  adjust.thresholds = TRUE )  
{  
  ...  
}
```

Protein Recognition: 28 Separate Models

```
# retrain best models on entire training set -----
```

```
train.existing.models.full = function(  
  model.folder.name,  
  structures = protein.names,  
  resolution = c(512,512),  
  training.bias = default.training.bias,  
  batch_size = 20,  
  steps_per_epoch = 20,  
  epochs = 80,  
  save.history.name = "history.Rda")  
{  
  ...  
}
```

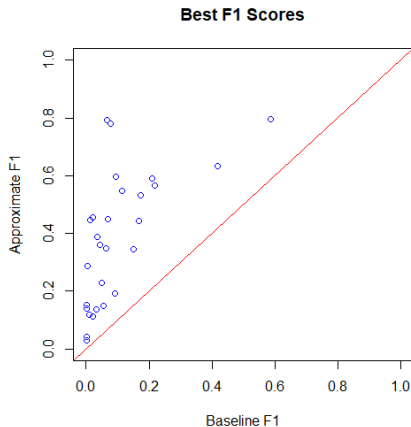

Protein Recognition: 28 Separate Models

```
# make a submission file to kaggle -----
```

```
kaggle.submission = function(  
  model.folder.name,  
  resolution = c(512,512),  
  adjust.thresholds = TRUE,  
  save.name = "kaggle.submission.csv")  
{  
  ...  
}
```



Protein Recognition: Results

Kaggle score: .24 Macro F_1 .





Protein Recognition: Results

Good:





Protein Recognition: Results

Good:





Protein Recognition: Results

Good:





Protein Recognition: Results

Good:





Protein Recognition: Results

Not good:



Protein Recognition: Results

Not good:



House Prices: Overview


- To predict: Sale price of a home
- 78 predictors (square footage, bedrooms, neighborhood, etc.)
- Training data: 1460 homes (price known)
- Test data: 1459 homes (price unknown)

House Prices: Metric


Root mean square logarithmic error:

$$RMSLE^2 = \frac{1}{N} \sum_{i=1}^N (\log(1 + \hat{y}_i) - \log(1 + y_i))^2$$


House Prices: Data Exploration



House Prices: Data Exploration



House Prices: Data Exploration



House Prices: Special Considerations

- Incomplete, redundant, and poorly formatted data
- Sale price is log-normally distributed
- Some predictors will be more informative than others
- Lots of models could possibly work here

House Prices: Data Cleanup


Steps taken to make the data usable:

- Data orthogonalization
- Level condensing
- Rank-deficiency correction
- Outlier removal

House Prices: Results

Learning Method	CV method	CV RMSLE	Test RMSLE	Notes
ordinary least squares	K = 10, B = 20	0.0988	0.13	all predictors used
least squares with subset selection	K = 10, B = 20	0.0908	0.1264	Forward subset selection identified 199 possible predictor subsets; CV was performed on each; lowest CV achieved with 83 predictors; this was submitted for test RMSLE
principal components regression	LOOCV	0.0989	0.1294	lowest CV achieved with 195 out of 199 principal components; this was submitted for test RMSLE
ridge regression	K = 10, B = 1	0.098	0.12419	Examined 150 values of λ from 10^{10} to 10^{-5} ; lowest CV achieved at $\lambda = .01$; this was submitted for test RMSLE
LASSO	K = 10, B = 1	0.09587	0.12641	Lowest CV achieved at $\lambda = .0001913724$; this was submitted for test RMSLE
regression tree	K = 10, B = 5	0.203		Didn't bother submitting, CV error too high
regression bag	OOB	0.1261	0.1528	500 trees grown (OOB error optimized long before that)
random forest	OOB	0.1234	0.1506	9 predictors at a time considered, 500 trees grown (OOB error optimized long before that)

House Prices: A Cool Graph



Thank you!